

Operating Systems

Introduction

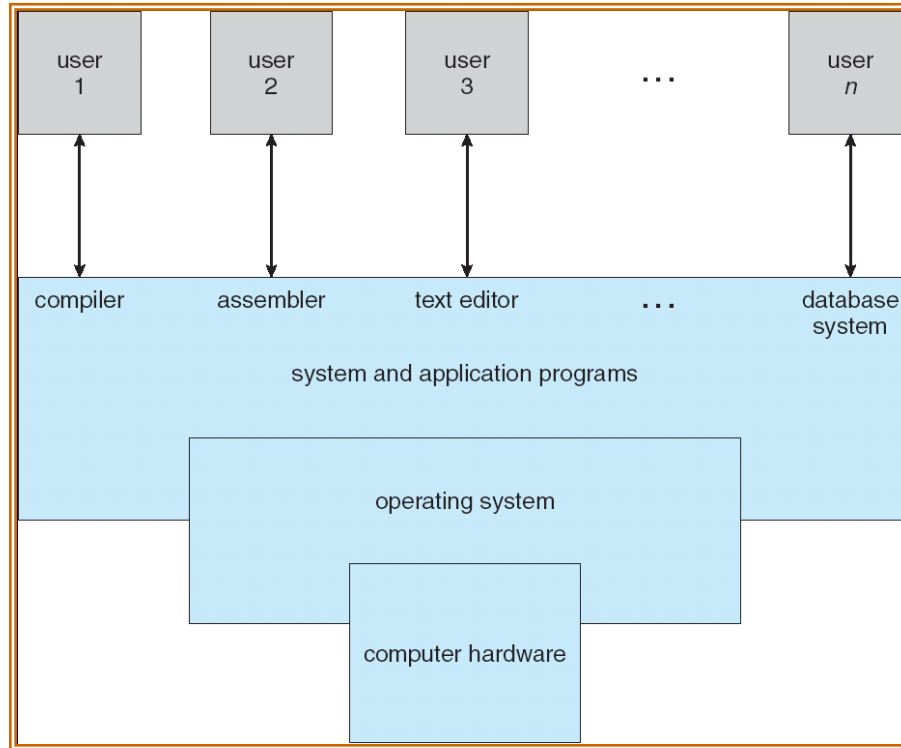
Computing Devices Everywhere



How do we tame complexity?

- Every piece of computer hardware different
 - Different CPU
 - Pentium, PowerPC, MIPS, ...
 - Different amounts of memory, disk, ...
 - Different types of devices
 - Mice, Keyboards, Sensors, Cameras, Fingerprint readers
 - Different networking environment
 - Cable, DSL, Wireless, Firewalls,...
- Questions:
 - Does the programmer need to write a single program that performs many independent activities?
 - Does every program have to be altered for every piece of hardware?
 - Does a faulty program crash everything?
 - Does every program have access to all hardware?

Four Components of a Computer System



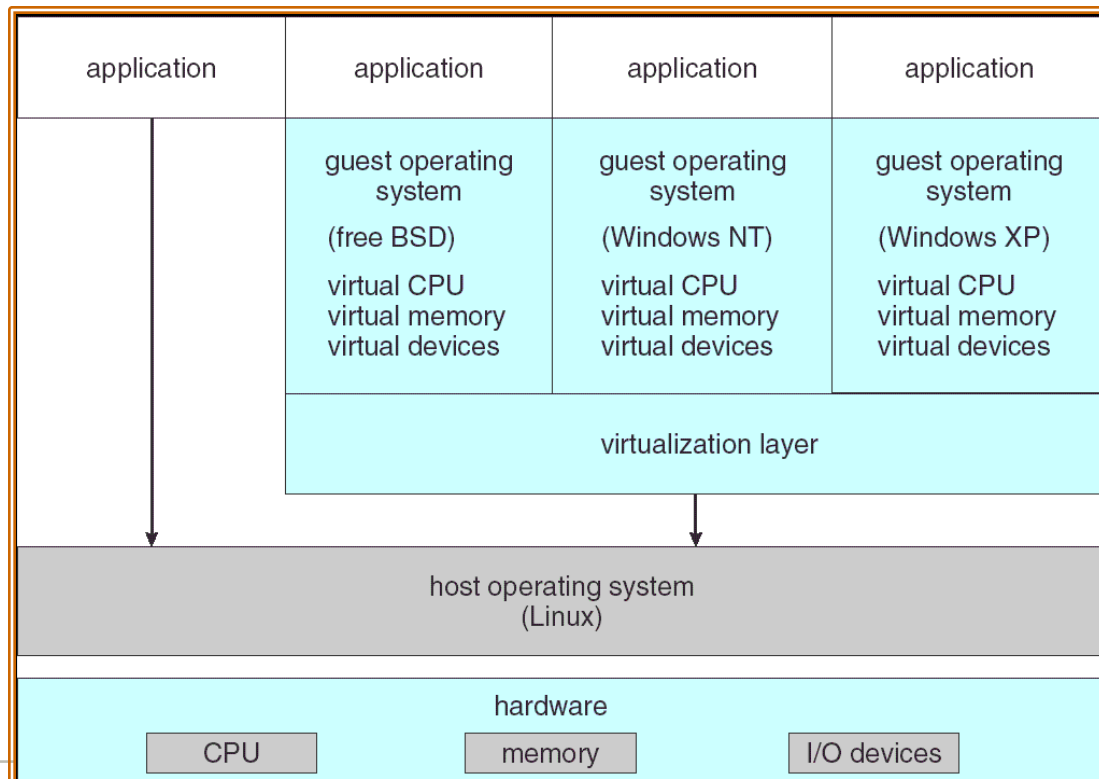
- An operating system is the interface between the user and the hardware.
- Definition: An operating system implements a virtual machine that is (hopefully) **easier and safer to program and use** than the raw hardware.

OS as Virtual Machines

- Software emulation of an abstract machine
 - Make it look like hardware has features you want
 - Programs from one hardware & OS work on another one
- Fault Isolation
 - Processes unable to directly impact other processes
 - Bugs cannot crash whole machine
- Protection and Portability
- Programming simplicity
 - Each process thinks it has all memory/CPU time
 - Each process thinks it owns all devices
 - Different Devices appear to have same interface
 - Device Interfaces more powerful than raw hardware
 - Bitmapped display \Rightarrow windowing system
 - Ethernet card \Rightarrow reliable, ordered, networking (TCP/IP)

Virtual Machines: Layers of OSs

- Useful for OS development
 - When OS crashes, restricted to one VM
 - Can aid testing programs on other OSs



Operating System: System View

- OS is a **resource allocator**
 - Manages all resources, including file system, virtual memory, networking, CPU etc.
 - Decides between conflicting requests for efficient and fair resource use
- OS is a **control program**
 - Controls execution of programs to prevent errors and improper use of the computer
- The OS design goal is to make the machine convenient to use (a software engineering problem) and efficient (a system and engineering problem).

Operating System Definition

- No universally accepted definition
- “Everything a vendor ships when you order an operating system” is good approximation
 - But varies wildly
- “The one program running at all times on the computer” is the **kernel**. Everything else is either a **system program** (ships with the operating system) or an **application program**

System Boot

- Operating system must be made available to hardware so hardware can start it
 - When power initialized on system, execution starts at a fixed memory location
 - Small piece of code – **bootstrap loader**, stored in Firmware **ROM** locates the kernel, loads it into memory, and starts it
 - Sometimes two-step process where **boot block** at fixed location loaded by ROM code, which loads bootstrap program from disk
 - Common bootstrap loader, **GRUB**, allows selection of kernel from multiple disks, versions, kernel options
- Kernel loads and system is then **running**

Why Study Operating Systems?

- The operating system makes the computer work - it is a key abstraction layer for applications
 - How to get the OS to give users an illusion of infinite memory, CPUs, resources, world wide computing, etc.
- System Design: How to make tradeoffs between
 - performance and the convenience of OS abstractions,
 - performance and the simplicity of OS design, and
 - putting functionality in hardware or software.
- As systems change the OS must adapt (e.g., new hardware, software).

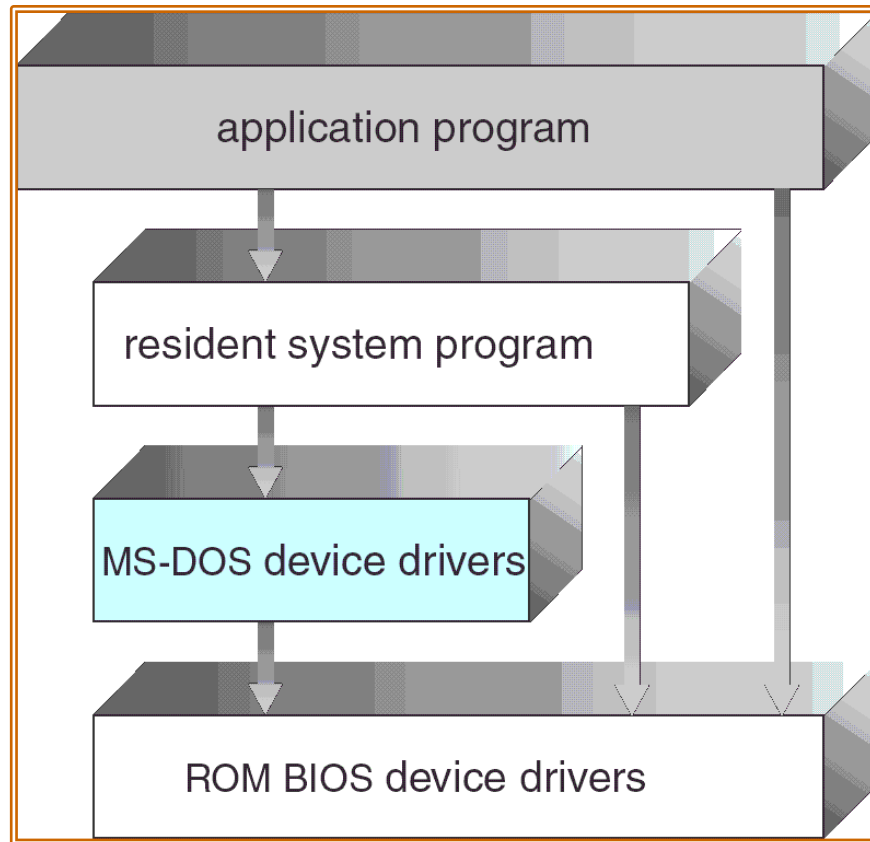
What if we didn't have an Operating System?

- Source Code \Rightarrow Compiler \Rightarrow Object Code \Rightarrow Hardware
- How do you get object code onto the hardware?
- How do you print out the answer?
- Once upon a time, had to toggle in program in binary and read out answer from LED's!

Simple OS: What if only one application?

- Examples:
 - Very early computers
 - Early PCs
 - Embedded controllers (elevators, cars, etc)
- OS becomes just a library of standard services
 - Standard device drivers
 - Interrupt handlers
 - Math libraries

MS-DOS Layer Structure



More complex OS: Multiple Apps

- Full Coordination and Protection
 - Manage interactions between different users
 - Multiple programs running simultaneously
 - Multiplex and protect Hardware Resources
 - CPU, Memory, I/O devices like disks, printers, etc
- Facilitator
 - Still provides Standard libraries, facilities
- Next we look a little at the history of OS

Moore's Law Change Drives OS Change

	1981	2006	Factor
CPU MHz, Cycles/inst	10 3—10	3200x4 0.25—0.5	1,280 6—40
DRAM capacity	128KB	4GB	32,768
Disk capacity	10MB	1TB	100,000
Net bandwidth	9600 b/s	1 Gb/s	110,000
# addr bits	16	32	2
#users/machine	10s	≤ 1	≤ 0.1
Price	\$25,000	\$4,000	0.2


Typical academic computer 1981 vs. 2006

Moore's law effects

- Nothing like this in any other area of business
- Transportation in over 200 years:
 - 2 orders of magnitude from horseback @10mph to Concorde @1000mph
 - Computers do this every decade!
- What does this mean for us?
 - Techniques have to vary over time to adapt to changing tradeoffs
- I place a lot more emphasis on principles
 - The key concepts underlying computer systems
 - Less emphasis on facts that are likely to change over the next few years...
- Let's examine the way changes in \$/MIPS has radically changed how OS's work

Dawn of time

ENIAC: (1945—1955)



- “The machine designed by Drs. Eckert and Mauchly was a monstrosity. When it was finished, the ENIAC filled an entire room, weighed thirty tons, and consumed two hundred kilowatts of power.”
- <http://ei.cs.vt.edu/~history/ENIAC.Richey.HTML>

History Phase 1 (1948—1970)

Hardware Expensive, Humans Cheap

- When computers cost millions of \$'s, optimize for more efficient use of the hardware!
 - Lack of interaction between user and computer
- **User at console**: one user at a time
- **Batch monitor**: load program, run, print
- Optimize to better use hardware
 - When user thinking at console, computer idle \Rightarrow BAD!
 - Feed computer batches and make users wait
- *No protection*: what if batch program has bugs?


History Phase 1 ½ (late 60s/early 70s)

- **Data channels, Interrupts:** overlap I/O and compute
 - DMA – Direct Memory Access for I/O devices
 - I/O can be completed asynchronously
- **Multiprogramming:** several programs run simultaneously
 - Small jobs not delayed by large jobs
 - More overlap between I/O and CPU
 - Need memory protection between programs and/or OS
- **Complexity gets out of hand:**
 - Multics: announced in 1963, ran in 1969
 - 1777 people “contributed to Multics” (30-40 core dev)
 - Turing award lecture from Fernando Corbató (key researcher): “On building systems that will fail”
 - OS 360: released with 1000 known bugs (APARs)
 - “Anomalous Program Activity Report”
- **OS finally becomes an important science:**
 - How to deal with complexity???
 - UNIX based on Multics, but vastly simplified

History Phase 2 (1970 – 1985)

Hardware Cheaper, Humans Expensive

- Computers available for tens of thousands of dollars instead of millions
- OS Technology maturing/stabilizing
- *Interactive timesharing*:
 - Use cheap terminals (~\$1000) to let multiple users interact with the system at the same time
 - Sacrifice CPU time to get better response time
 - Users do debugging, editing, and email online
- **Problem: Thrashing**
 - Performance very non-linear response with load
 - Thrashing caused by many factors including
 - Swapping, queuing



History Phase 3 (1981—)

Hardware Very Cheap, Humans Very Expensive

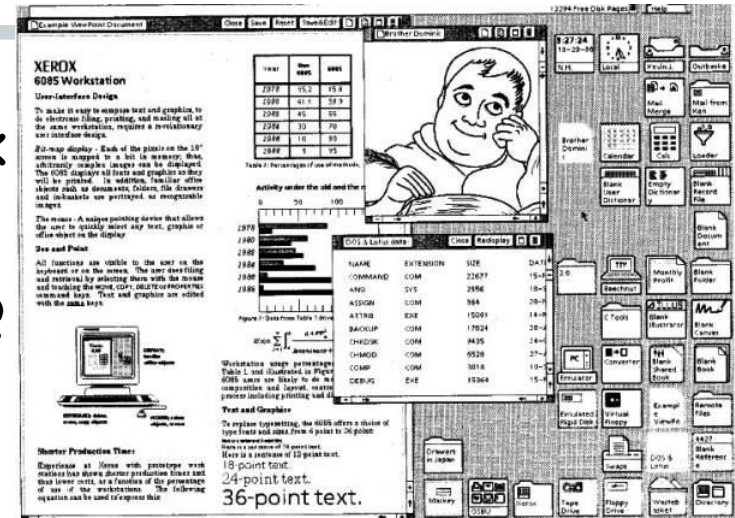
- Computer costs \$1K, Programmer costs \$100K/year
 - If you can make someone 1% more efficient by giving them a computer, it's worth it!
 - Use computers to make people more efficient
- **Personal computing:**
 - Computers cheap, so give everyone a PC
- **Limited Hardware Resources Initially:**
 - OS becomes a subroutine library
 - One application at a time (MSDOS, CP/M, ...)
- **Eventually PCs become powerful:**
 - OS regains all the complexity of a “big” OS
 - multiprogramming, memory protection, etc (Windows NT, OS/2)
- Question: As hardware gets cheaper does need for OS go away?

History Phase 3 (con't)

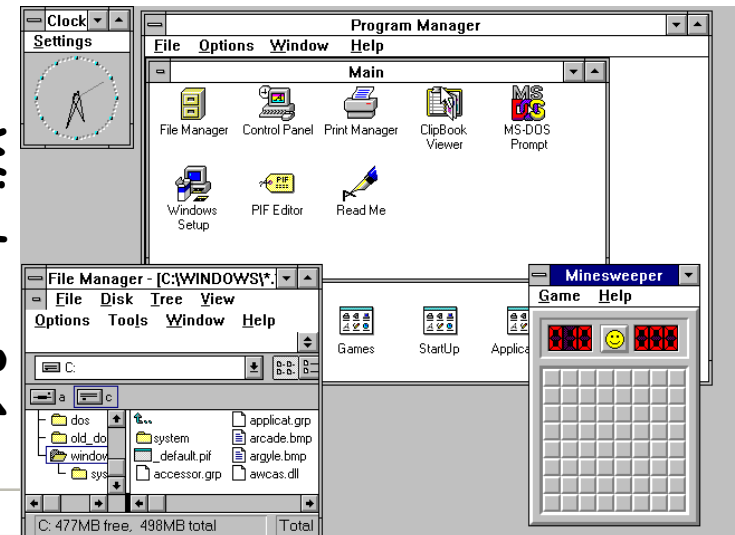
Graphical User Interfaces

- Xerox Star: 1981
 - Originally a research project (Alto)
 - First “mice”, “windows”
- Apple Lisa/Machintosh: 1984
 - “Look and Feel” suit 1988
- Microsoft Windows:
 - Win 1.0 (1985)
 - Win 3.1 (1990)
 - Win 95 (1995)
 - Win NT (1993)
 - Win 2000 (2000)
 - Win XP (2001)
 - Win Vista (2007)
 - Win 7 (2009)

Xerox Star




Windows 3.1



History Phase 4 (1989—): Distributed Systems

- Networking (Local Area Networking)
 - Different machines share resources
 - Printers, File Servers, Web Servers
 - Client – Server Model
- Services
 - Computing
 - File Storage



History Phase 5 (1995—): Mobile Systems


- Ubiquitous Mobile Devices
 - Laptops, PDAs, smart phones
 - Small, portable, and inexpensive
 - Recently twice as many smart phones as PDAs
 - Many computers/person!
 - Limited capabilities (memory, CPU, power, etc...)
- Wireless/Wide Area Networking
 - Leveraging the infrastructure
 - Huge distributed pool of resources extend devices
 - Traditional computers split into pieces. Wireless keyboards/mice, CPU distributed, storage remote
- Peer-to-peer systems
 - Many devices with equal responsibilities work together
 - Components of “Operating System” spread across globe

History of OS: Summary

- Change is continuous and OSs should adapt
 - Not: look how stupid batch processing was
 - But: Made sense at the time
- Situation today is much like the late 60s [\[poll\]](#)
 - Small OS: 100K lines
 - Large OS: >10M lines (5M for the browser!)
 - 100-1000 people-years
- Complexity still reigns
 - NT under development from early 90's to late 90's
 - Never worked very well
 - Windows XP has 45M lines of code,
 - Windows Vista delayed many times
 - Latest release date of 2005, 2006, 2007


Computer System Organization

- Computer-system operation
 - One or more CPUs, device controllers connect through common bus providing access to shared memory
 - Concurrent execution of CPUs and devices competing for memory cycles




Computer-System Architecture

- Many systems use a single general-purpose processor (PDAs through mainframes)
 - Most systems have special-purpose processors as well.
- Multiprocessors systems growing in use and importance
 1. [Asymmetric Multiprocessing](#)
 2. [Symmetric Multiprocessing](#)




A Dual-Core Design



Storage Hierarchy

- Storage systems organized in hierarchy
 - Speed
 - Cost
 - Volatility
- Main memory – only large storage media that the CPU can access directly
- Magnetic (hard) disks – rigid metal or glass platters covered with magnetic recording material
 - Disk surface is logically divided into **tracks**, which are subdivided into **sectors**
 - The **disk controller** determines the logical interaction between the device and the computer
- **Caching** – copying information into faster storage system; main memory can be viewed as a last *cache* for secondary storage

Storage-Device Hierarchy



Performance of Various Levels of Storage

- Movement between levels of storage hierarchy can be explicit or implicit

Level	1	2	3	4
Name	registers	cache	main memory	disk storage
Typical size	< 1 KB	> 16 MB	> 16 GB	> 100 GB
Implementation technology	custom memory with multiple ports, CMOS	on-chip or off-chip CMOS SRAM	CMOS DRAM	magnetic disk
Access time (ns)	0.25 – 0.5	0.5 – 25	80 – 250	5,000.000
Bandwidth (MB/sec)	20,000 – 100,000	5000 – 10,000	1000 – 5000	20 – 150
Managed by	compiler	hardware	operating system	operating system
Backed by	cache	main memory	disk	CD or tape

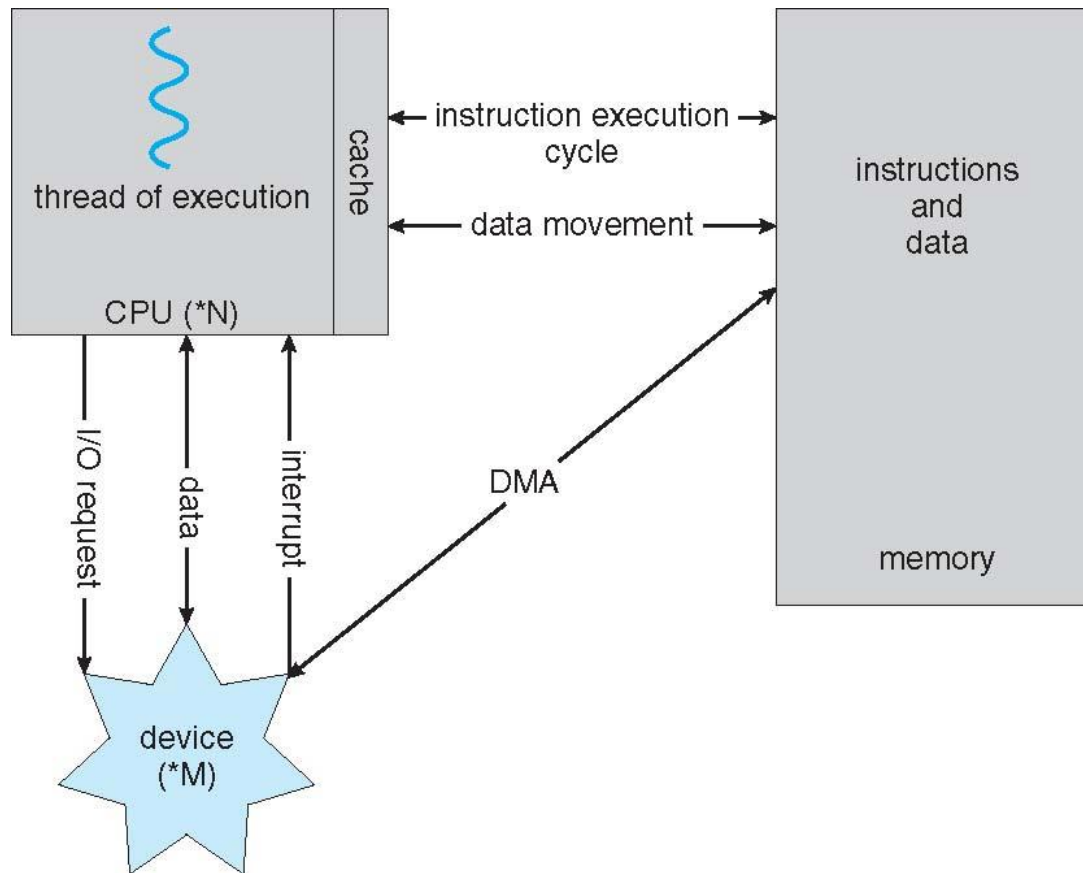
Caching

- Important principle, performed at many levels in a computer (in hardware, operating system, software)
- Information in use copied from slower to faster storage temporarily
- Faster storage (cache) checked first to determine if information is there
 - If it is, information used directly from the cache (fast)
 - If not, data copied to cache and used there
- Cache smaller than storage being cached
 - Cache management important design problem
 - Cache size and replacement policy

Computer-System Operation

- I/O devices and the CPU can execute concurrently
- Each device controller is in charge of a particular device type
- Each device controller has a local buffer
- CPU moves data from/to main memory to/from local buffers
- I/O is from the device to local buffer of controller
- Device controller informs CPU that it has finished its operation by causing an *interrupt*

How a Modern Computer Works



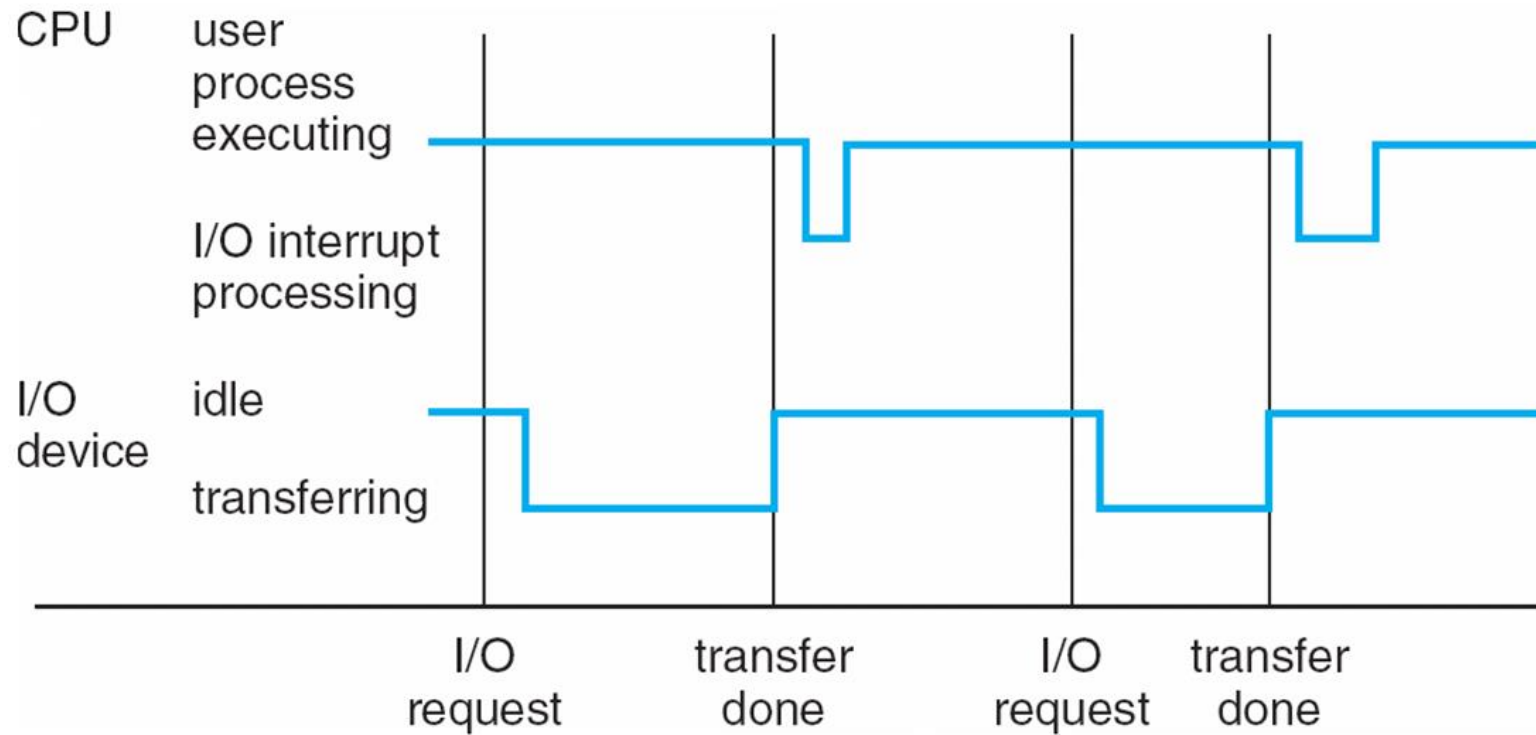
Interrupt Driven O.S.

- Most modern operating systems are interrupt driven.
 - Start-up: 1) Load O.S. (kernel) and start it running.
 - 2) O.S. waits for an event (an interrupt).
- Definition of interrupt:
 - An **interrupt** is a method to ensure that the CPU takes note of an event.
- Types of interrupt:
 - Hardware interrupts (e.g. from an I/O device)
 - Software interrupt (from an executing process)
 - Also called a **trap** or an **exception**
 - Generated to signal error (e.g. divide by zero)
 - or to request service from OS (system calls for I/O requests).

Interrupt Handling

- The operating system preserves the state of the CPU by storing registers and the program counter.
- The O.S. determines which type of interrupt has occurred.
 - Separate segments of code determine what action should be taken for each type of interrupt.
- The OS executes the sequence of commands associated with the given interrupt.
- The OS recovers the stored information from the original process and continues execution.

Interrupt Timeline



Summary

- Operating systems provide a virtual machine abstraction to handle diverse hardware
- Rapid Change in Hardware Leads to changing OS
 - Batch \Rightarrow Multiprogramming \Rightarrow Timeshare \Rightarrow Graphical UI \Rightarrow Ubiquitous Devices \Rightarrow Cyberspace/Metaverse/??
- OS features migrated from mainframes \Rightarrow PCs \Rightarrow Mobile

Summary of Operating System Principles

- OS as juggler: providing the illusion of a dedicated machine with infinite memory and CPU.
- OS as government: protecting users from each other, allocating resources efficiently and fairly, and providing secure and safe communication.
- OS as complex system: keeping OS design and implementation as simple as possible is the key to getting the OS to work.
- OS as history teacher: learning from past to predict the future, i.e., OS design tradeoffs change with technology.